Suggested Architecture

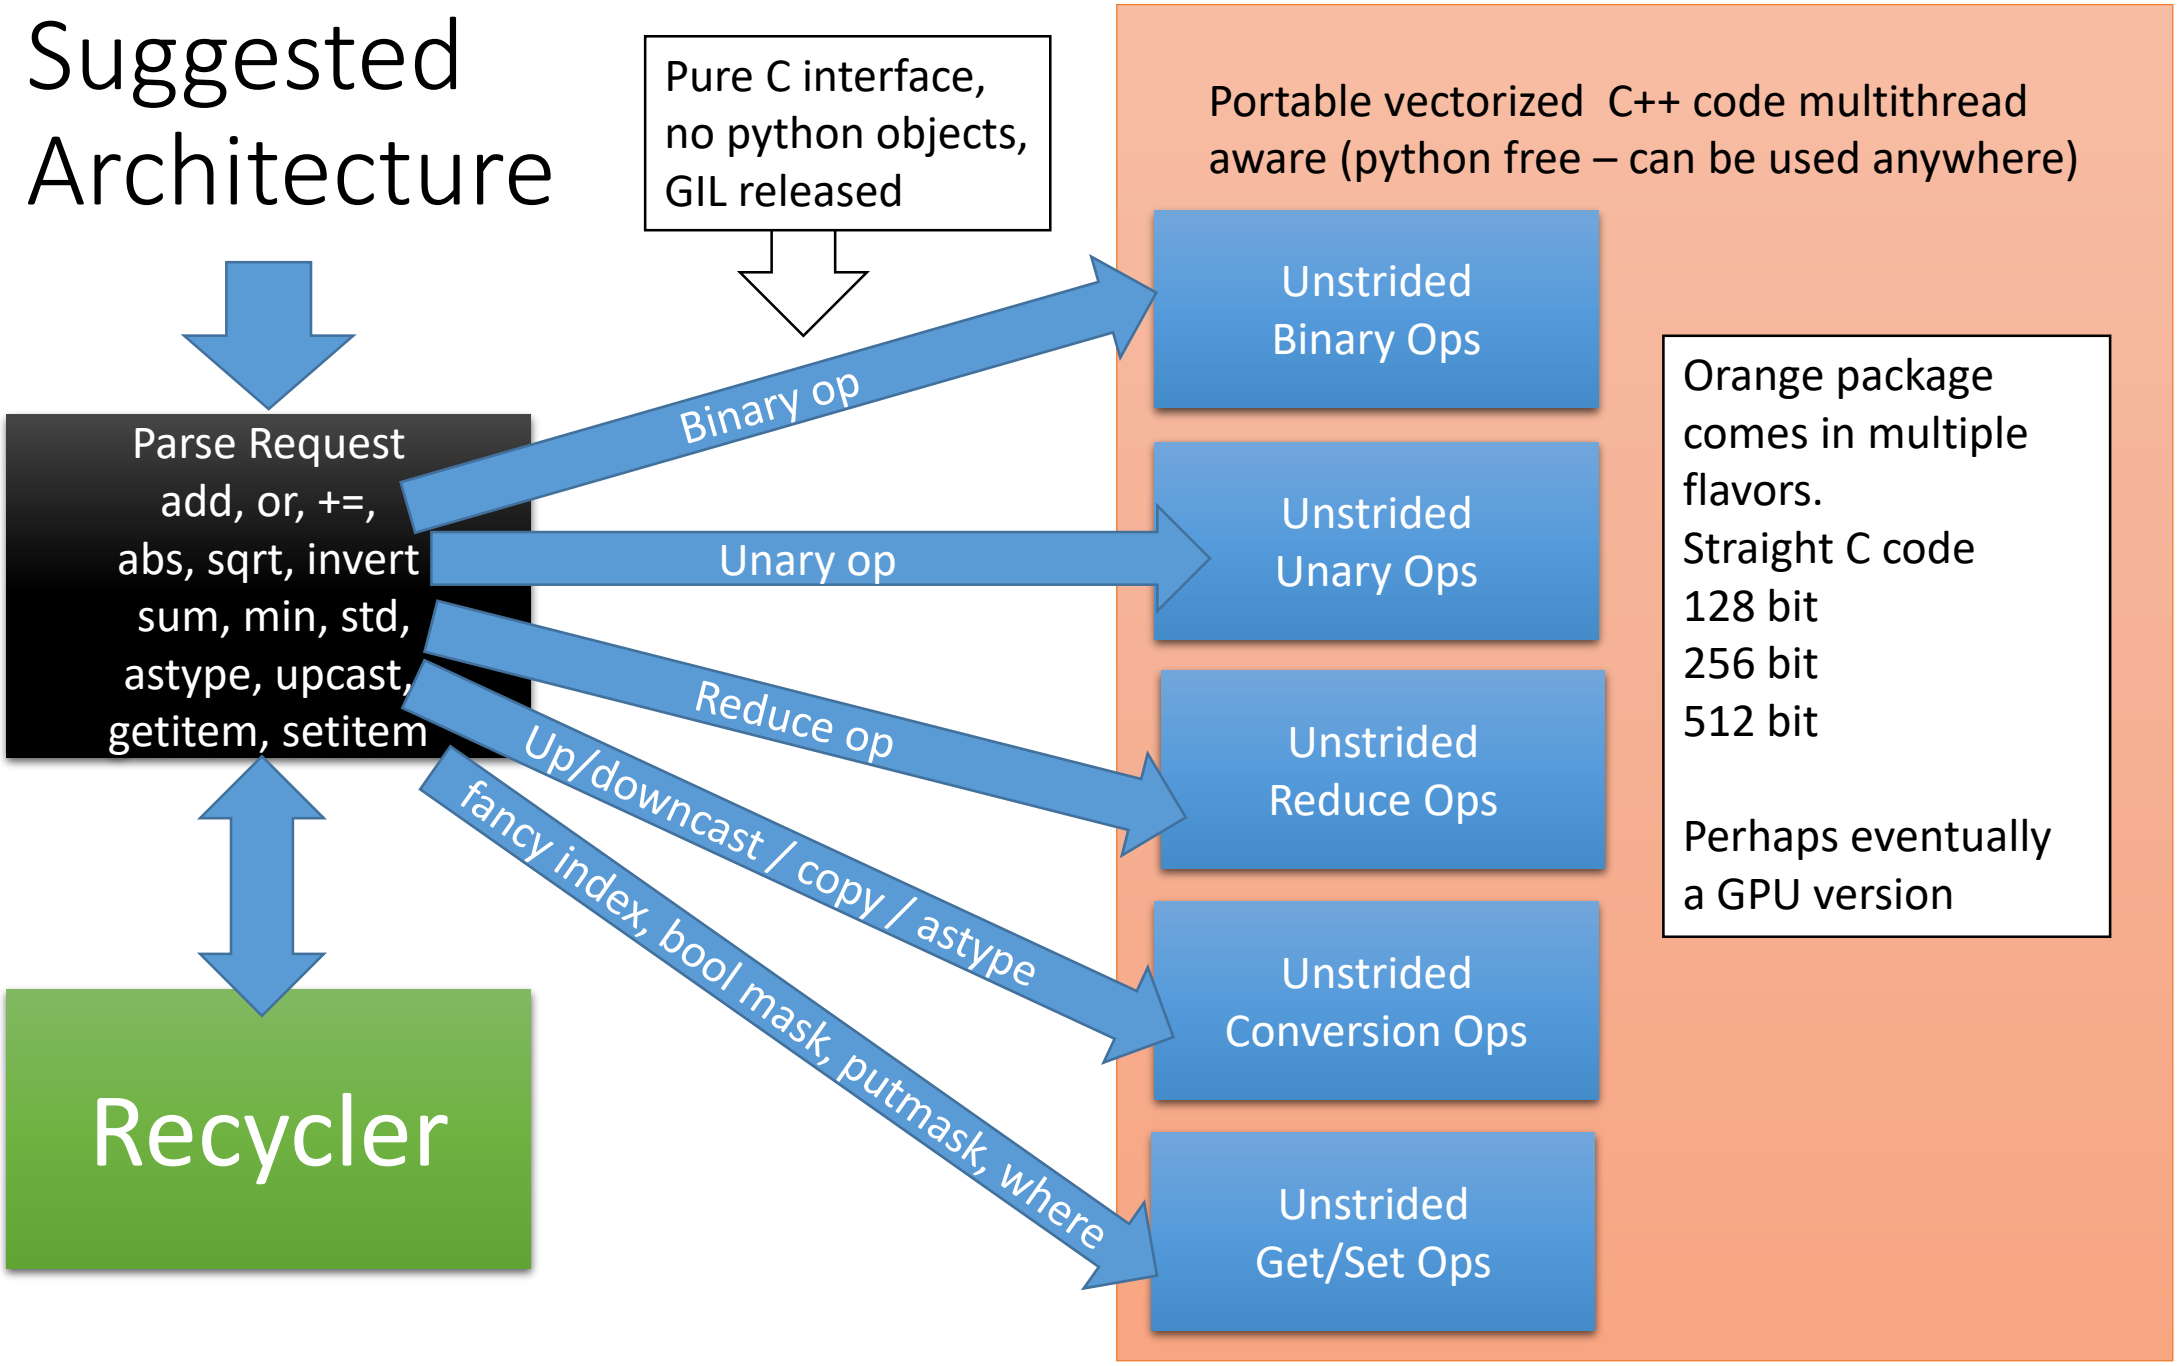

Suggested approach

- Recycling – off by default if large memory supported, else default on, universal to all platforms, can tune (see write up)
- Threading – off by default, universal to all platforms, can turn on, assign numa affinity, assign max count, etc.
- Vectorization – on by default for computers that support it

Vectorization is chip dependent

- 128 / 256/ 512 flavors
  - Start with 256. If computer only supports 128 – do they really need high performance since that chip is 8+ years old? Not worth effort initially.
- Requires plug in package so others can contribute – for instance, once a user sees how 256 bit package is written, the 512 package is derivable (and so is 128 or some future chip).
- Newer chips have more features that can be exploited (and thus the pluggable package approach).

First step

- Write the recycler first otherwise – you will not be able to tell when CPU is hitting max speeds due to page faults.
- Write threading second

ds.groupby(<ONE OR MORE KEYS>, kwargs).OPERATION(<ONE OR MORE INPUTS>, kwargs)

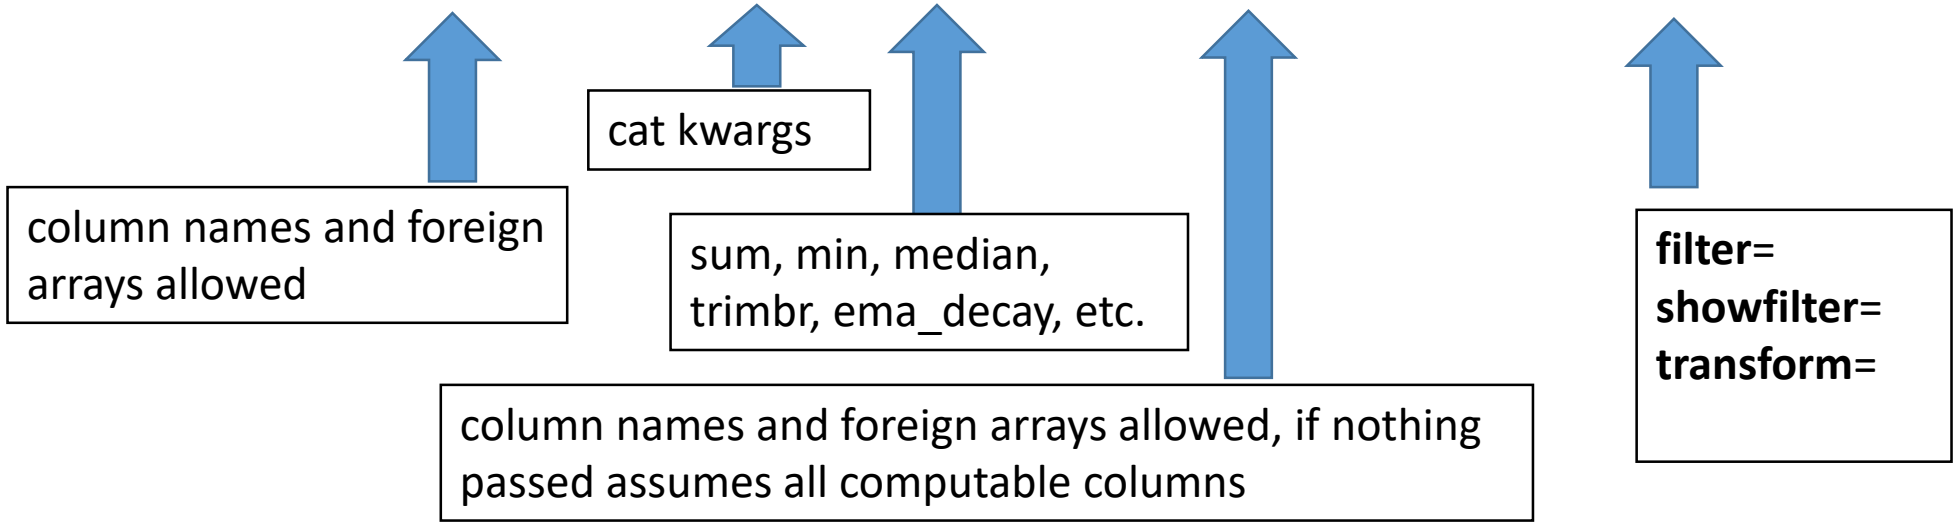

ds.groupby(<ONE OR MORE KEYS>).agg(<ONE OR MORE OPS>, <ONE OR MORE INPUTS>, kwargs)

ds.groupby(<ONE OR MORE KEYS>).apply_reduce(<ONE OR MORE OPS>, <ONE OR MORE INPUTS>, kwargs)

Groupby shorthand methods

ds.gb: does pandas style – first occurrence binning, but display sorted

ds.gbu: 'u' = unordered. First occurrence ordering and binning. Bins displayed in first occurrence.

ds.gbs: 's' = sorted. All keys sorted (lexsort). Bins are sorted and displayed that way. (lex=True)

Categorical Form then matches identically:

cat.OPERATION(<ONE OR MORE INPUTS>, kwargs)

cat.agg(<ONE OR MORE OPS>, <ONE OR MORE INPUTS>, kwargs)

cat.apply_reduce(<ONE OR MORE OPS>, <ONE OR MORE INPUTS>, kwargs)

Multiple ways to produce a new column ‘F’ which can be fed back in